Workflow Assessment: AI Integration & Risk Evaluation

Apex Strategy Group | Government welfare-policy engagement | Internal risk and workflow redesign

Scope note: *The assignment instructions identify Task 4 as Filling Documentation Gaps but do not include the original five-task case study or downloadable Word template. Tasks 1, 2, 3 and 5 below are clearly labeled working reconstructions drawn from the described failures; their exact original wording should be checked against the course case study before submission.*

Section 1 | Task Inventory and AI Involvement

Scoring: 1 = low; 5 = high. Accountability measures the consequence of an incorrect output; sensitivity measures confidentiality and public impact; complexity measures analytical judgment and evidence requirements.

Task (working reconstruction)	Account.	Sens.	Complex.	AI involvement	Original workflow gap
1. Locate policy sources and references	4	3	3	AI-assisted	AI-suggested references entered the research trail without existence and provenance checks.
2. Extract and analyze welfare statistics	5	4	5	AI-assisted	Generated figures were not reconciled to original statistical tables, definitions, or reporting periods.
3. Summarize policy evidence and draft analysis	5	4	4	AI-assisted	The team treated fluent AI summaries as verified interpretations without reading the underlying documents.
4. Fill documentation gaps with missing data	5	5	5	Human-led	When evidence was unavailable, AI invented missing source material and presented it as factual support.
5. Assemble and quality-check final report	5	4	4	Human-led	No independent human citation/data audit or accountable release approval blocked the defective submission.

Section 2 | Redesigned Action Plan (Three Priorities)

Priority 1 - Task 2: Extract and analyze welfare statistics

AI role: Produce proposed variable labels, comparison questions, and reproducible calculation code using only analyst-approved, de-identified extracts; never supply unsourced numeric values.

Source of truth: The relevant government agency's official statistical release, accompanying downloadable tables, metadata, and documented reporting period.

Verification: A human analyst opens the official table, matches each report value to its original cell and denominator, independently recalculates transformations, and logs the URL/document title, table ID, publication date, and reviewer initials.

Safeguard: The report build rejects statistics without an evidence-register entry and second-person signoff; confidential client extracts remain in approved enterprise systems, not public AI tools.

Priority 2 - Task 4: Filling Documentation Gaps

AI role: Only identify missing claims, draft neutral clarification questions, and create a clearly labeled evidence-gap log; AI must not generate surrogate citations, facts, or fabricated source passages.

Source of truth: The original client-approved evidence package, primary policy document cited in the engagement brief, and the agency's official publication archive.

Verification: The lead consultant manually checks source availability, records unsuccessful searches, requests the missing document from its owner, and either verifies the original evidence or states the limitation explicitly.

Safeguard: Mandatory STOP gate: no unsupported claim, placeholder, or synthetic citation can enter the final report; unresolved gaps are excluded or disclosed and escalated to the engagement manager.

Priority 3 - Task 5: Final report assembly and quality assurance

AI role: Check style consistency and flag uncited sentences or apparent mismatches against an approved evidence register; AI may propose wording but cannot authorize publication.

Source of truth: The signed engagement brief, approved source/evidence register, original government publications, and final analyst-validated calculation workbook.

Verification: An independent reviewer opens every cited primary document, confirms title/author/date/page and that each citation supports the precise claim; reconciles all tables and verifies conclusions against the underlying analysis.

Safeguard: Two-person release checklist and version-controlled signoff by research lead and engagement manager; any failed check blocks client delivery.

Section 3 | Highest-Risk Activity: Task 4

Core risk: Using AI to "fill" missing source data creates fabricated evidence with the appearance of authority. Because downstream policy advice could rely on unsupported statements, this is not an acceptable gap-filling method even when the generated text seems plausible.

Primary Microsoft Responsible AI principle violated: Reliability & Safety. Outputs that invent unverified data are neither dependable nor safe for consequential public-policy decisions. **Related violations:** Transparency (synthetic material was not labeled as such) and Accountability (no responsible human verified or rejected the content).

Decision: Human-led only for evidence resolution. Record the missing item, request an authoritative source, disclose any unresolved limitation, and prohibit fabricated content from entering a deliverable.

Section 4 | Personal Reflection and Commitment

In my future AI-assisted research and consulting work, I will use a source-verification gate before accepting any AI-generated statistic, citation, or factual claim into a client-facing document. For every material statement, I will open the original primary source, confirm that it exists, compare the specific number or wording with the document, and record the source title, date, page or table, and verification status in an evidence register. If I cannot establish the source, I will label the claim as unverified, remove it from the draft, or escalate the evidence gap rather than asking AI to invent a plausible replacement. This habit is especially important in healthcare analytics and public-policy work, where unsupported claims can distort decisions affecting real people. I will still use AI for bounded tasks such as brainstorming search terms, identifying missing references, and editing the clarity of verified prose. However, I will reserve factual approval and final signoff for accountable human reviewers. A traceable evidence trail will make my work more reproducible, auditable, and trustworthy while retaining the efficiency benefits of responsible AI assistance.

Reflection word count: 177 words.
Completion check: 5 tasks scored | 3 detailed redesigns | Task 4 prohibited from fabrication | AI drafts separated from verified claims | personal commitment within 150-200 words.